

# Balanced BST and AVL Trees



***Carlos Moreno***

cmoreno@uwaterloo.ca

EIT-4103



**<https://ece.uwaterloo.ca/~cmoreno/ece250>**

These slides, the course material, and course web site are based on work by Douglas W. Harder

# Balanced BST and AVL Trees

Standard reminder to set phones to  
silent/vibrate mode, please!



# Balanced BST and AVL Trees

- Previously, on ECE-250...
  - We talked about Binary Search Trees (BST) and their applicability to searches.
  - We also wanted to use them to maintain a list of ordered values (as in, adding and removing elements).
    - And of course, greedy as we are, never happy enough with what we have, we *also* wanted to do all of this efficiently!!
  - We discussed that this required the height of the tree to be  $\Theta(\log n)$ , which requires a *balanced* tree.

# Balanced BST and AVL Trees

- And now...
  - We'll expand on this issue of Balanced Binary Search Trees.
    - We'll look into two “obvious” types of balance (weight and height).
  - We'll discuss AVL trees — one of the commonly used techniques to efficiently maintain a balanced BST.
    - We'll look into the techniques to perform all the operations and techniques to re-balance the tree.

# Balanced BST and AVL Trees

- Recall search on a binary search tree:
  - Data is not hierarchical, but rather linearly or totally ordered.
  - Values in the left sub-tree are all less than the value at the root, and values in the right sub-tree are all greater than the value at the root (we'll continue to assume that we have no duplicate values).
  - This allows us to do the analogous to a binary search — if the value we're searching is less than the current node, we continue searching through the left sub-tree.

# Balanced BST and AVL Trees

- Recall search on a binary search tree:
  - The idea being: if each sub-tree always has half as many elements, then we're doing essentially the same as in binary search — with each comparison, we discard half the remaining values.
    - This leads to logarithmic time in the search.

# Balanced BST and AVL Trees

- Recall search on a binary search tree:
  - An alternative way to look at this is: the search takes worst-case  $\Theta(h)$ , where  $h$  is the height of the tree
    - This tells us that a tree where the left and right sub-trees of every node have the same number of nodes has logarithmic height with respect to the number of nodes!

## Balanced BST and AVL Trees

- We notice that a tree where we insert and remove elements can not have strict balance as described so far:
  - Our definition has to accommodate a  $\pm 1$  difference (*why?*)



## Balanced BST and AVL Trees

- We notice that a tree where we insert and remove elements can not have strict balance as described so far:
  - Our definition has to accommodate a  $\pm 1$  difference (*why?*)
  - As a simple counter-example: suppose we have a perfect tree, and add one element. How could both sub-trees have the exact same number of elements?

## Balanced BST and AVL Trees

- A different type of balance, though leading to the same outcome (logarithmic height with respect to the number of nodes) is height balance:
  - For every node in the tree, the left and right sub-trees have the same height,  $\pm 1$

## Balanced BST and AVL Trees

- We'll now look into AVL trees (named after its two inventors, Adelson-Velskii and Landis), a type of height-balanced binary search trees.

## Balanced BST and AVL Trees

- We'll now look into AVL trees (named after its two inventors, Adelson-Velskii and Landis), a type of height-balanced binary search trees.
- A binary search tree is an AVL tree if:
  - The difference in heights between left and right sub-trees is at most 1, and
  - Both sub-trees are AVL trees


## Balanced BST and AVL Trees

- The principle of operation is remarkably simple; let's look at this “prototypical” trick to maintain balance:
  - Let's add 3, 2, 1 (in that order) to a binary search tree:

③


## Balanced BST and AVL Trees

- The principle of operation is remarkably simple; let's look at this “prototypical” trick to maintain balance:
  - Let's add 3, 2, 1 (in that order) to a binary search tree:




## Balanced BST and AVL Trees

- Inserting 1 causes the tree to become imbalanced at node 3 (left sub-tree has height 1, and right sub-tree has height ... ?? you tell me?)



## Balanced BST and AVL Trees


- Inserting 1 causes the tree to become imbalanced at node 3 (left sub-tree has height 1, and right sub-tree has height ... ?? you tell me?) — we recall from the first class on trees, that by convention, an empty tree has height  $-1$






# Balanced BST and AVL Trees

- So, we rotate it towards the right:



# Balanced BST and AVL Trees

- So, we rotate it towards the right:



## Balanced BST and AVL Trees


- If we had inserted the sequence 3, 2, 1, the situation would have been essentially the same; the right sub-tree would be the deeper one in that case, so we rotate towards the left in that case.

## Balanced BST and AVL Trees

- Inserting the sequence 3, 1, 2, however, does make a significant difference — we end up with the following tree:


## Balanced BST and AVL Trees

- Inserting the sequence 3, 1, 2, however, does make a significant difference — we end up with the following tree:




## Balanced BST and AVL Trees

- Inserting the sequence 3, 1, 2, however, does make a significant difference — we end up with the following tree:
  - Clearly, we can't rotate to balance (right? *why?*)




# Balanced BST and AVL Trees

- However, we can do a double rotation:
  - First, rotate towards the left the sub-tree 1-2, and then we're in the exact same previous case (so we rotate towards the right)



# Balanced BST and AVL Trees

- However, we can do a double rotation:
  - First, rotate towards the left the sub-tree 1-2, and then we're in the exact same previous case (so we rotate towards the right)
  - The “net” effect is the following:






# Balanced BST and AVL Trees

- What about a more complicated situation?

# Balanced BST and AVL Trees


- We insert the value  $a$ , which makes it into the subtree  $B_L$ , without causing any imbalance in it.



(notice that a blue triangle denotes a tree with height  $h$ )


# Balanced BST and AVL Trees

- We insert the value  $a$ , which makes it into the subtree  $B_L$ , without causing any imbalance in it.




# Balanced BST and AVL Trees

- How do we fix the imbalance in  $f$ ? If we try to rotate, what do we do about  $B_R$ ?




# Balanced BST and AVL Trees

- The key detail is that we have to detach  $B_R$  from the tree (temporarily — we'll re-attach it after the rotation!)




# Balanced BST and AVL Trees

- The key detail is that we have to detach  $B_R$  from the tree (temporarily — we'll re-attach it after the rotation!)




# Balanced BST and AVL Trees

- We now see that  $f - b - B_L$  is the sequence that we have to rotate towards the right:



# Balanced BST and AVL Trees


- We rotate:






## Balanced BST and AVL Trees

- Then re-attach  $B_R$  at the obvious place (it is obvious, right?):




# Balanced BST and AVL Trees

- It really is the only “loose” branch where we can attach it — the question is: can we attach it there without breaking the tree's constraints?




## Balanced BST and AVL Trees

- $B_R$  can clearly go as the left sub-tree of  $f$  — it was originally below  $b$  which was part of  $f$ 's left sub-tree. And it can also be part of  $b$ 's right sub-tree (it was originally  $b$ 's right sub-tree!):



# Balanced BST and AVL Trees

- Another interesting detail is that the height of the tree with root  $b$  equals the original height of the tree with root  $f$  — this tells us that the insertion can not affect the balance of any ancestors!



## Balanced BST and AVL Trees


- In other words, this tells us that whenever an imbalance is created, we only need to address it at the deepest level where there is imbalance.

## Balanced BST and AVL Trees

- In other words, this tells us that whenever an imbalance is created, we only need to address it at the deepest level where there is imbalance.
  - Things were originally balanced everywhere.
  - Fixing something at a node with depth  $d$  fixes any imbalance in any node above it.
    - Since nothing else changed in the tree, it must be the case that fixing the imbalance at the deepest node where imbalance is present must be sufficient.


# Balanced BST and AVL Trees

- For example, the just added 54 causes imbalance at the root, at 55, and at 48:



# Balanced BST and AVL Trees

- However, it's pretty clear that a rotation around 48 fixes the imbalance at all points!






# Balanced BST and AVL Trees

- However, not everything is good news...


# Balanced BST and AVL Trees

- However, not everything is good news... the *uh-oh!* situation happens if the inserted element makes it to  $B_R$  instead of  $B_L$ :




# Balanced BST and AVL Trees

- If we try the same trick, we end up with a still imbalanced tree!




# Balanced BST and AVL Trees

- If we try the same trick, we end up with a still imbalanced tree!




# Balanced BST and AVL Trees

- Notice that what's really going on is that we have a situation similar to the second prototypical example, where we needed a double rotation!




# Balanced BST and AVL Trees

- We saw that we just need to transform the situation into a situation similar to the first case... (*how do we do that?*)




# Balanced BST and AVL Trees

- As we saw with the prototypical example, we first rotate towards the left around  $b$ . Then we're back in the simpler case!



# Balanced BST and AVL Trees


- We “zoom in” into the sub-tree  $B_R$ :






# Balanced BST and AVL Trees

- And insert a value (either  $c$  or  $e$ )




# Balanced BST and AVL Trees

- Going from  $f$  in the direction of the cause of the imbalance, we do a left – right; we first detach things:




# Balanced BST and AVL Trees

- Then readjust:




# Balanced BST and AVL Trees

- Do I need to make you guess where would  $D_L$  and  $D_R$  be attached? :-)




# Balanced BST and AVL Trees

- Both can go at the right of  $b$ , and both can go at the left of  $f$  — so, we place them according to the constraint due to  $d$ :




# Balanced BST and AVL Trees

- Both can go at the right of  $b$ , and both can go at the left of  $f$  — so, we place them according to the constraint due to  $d$ :




## Balanced BST and AVL Trees

- In either case —  $c$ , added below  $D_L$ , or  $e$ , added below  $D_R$  — we end up with a balanced tree:



# Balanced BST and AVL Trees


- An example: insertion of 46 causes imbalance at nodes 42, 55, and 36 — we address the imbalance at 42 (the deepest node):






# Balanced BST and AVL Trees

- First, rotate towards the right around 43 (detaching 46 and reattaching it as 48's left sub-tree):



# Balanced BST and AVL Trees

- Then rotate towards the left around 42:



## Balanced BST and AVL Trees

- We observe a piece of good news: these rotations take  $\Theta(1)$ , and the insertion takes  $\Theta(h) = \Theta(\log n)$ , so we're in good shape: insertion (including maintaining balance) takes  $\Theta(\log n)$
- We'll also see that removals, though a bit more complicated (and more inefficient), also take  $\Theta(\log n)$ .

# Summary

- During today's class:
  - Introduced the notion of balance in BSTs
  - Discussed height-balance (as opposed to weight-balance)
  - Looked into AVL trees:
    - Prototypical examples of re-balancing a tree that becomes unbalanced due to an insertion.
    - Discussed the single-rotation and double-rotation.